

Implement a function `related_words()` that returns a list of words (stemmed words are fine) that are semantically similar to the input. It should take two arguments:

1. `word` - the word of interest [string]
2. `n` - the number of similar words to return [int]

and return:

1. `related_words` - list of strings

Do this by computing the term-document matrix using tf-idf weighting, and measuring the cosine distance between word vectors.

- Use the tf-idf weightings suggested by Jurafsky, ch. 6.
- Use the news, editorial, and reviews corpora.

Put your function in a file titled `hw09_solution.py`.

Run `hw09_evaluate.py`, making sure that your solution file is on the Python path. If you're unsure whether the result is satisfactory, ask the instructor or TA.